

Scanning Agent – Document Scoring Metrics

The Scanning Agent evaluates the **structural quality** of an uploaded document using a universal rubric. The score does not determine whether the document is the correct document for a particular SDLC stage and does not evaluate factual accuracy.

Each document is evaluated on the following three criteria:

1. **Structural Clarity – 20 marks**

2. Checks whether the document has clear headings and sections.
3. Evaluates whether the content follows a logical flow.
4. Identifies documents that are presented as a wall of undifferentiated text without a navigable structure.

5. **Completeness – 20 marks**

6. Checks for blank fields and missing content.
7. Identifies placeholder text such as "TBD", "Lorem ipsum", or empty brackets.
8. Checks whether screenshots/images are being used in place of content that should be provided as actual text.

9. **Labeling Accuracy – 20 marks**

10. Checks whether the content under each heading actually corresponds to that heading.
11. For example, a section titled "Test Cases" should contain test cases rather than unrelated content.

Scoring

- Each criterion is scored from **0–20** using integer scores.
- **Maximum possible score: 60**
- The overall score is calculated as the sum of the three individual criterion scores.
- A document scoring **36/60 or below** falls below the current 60% threshold and is considered for reformation/improvement.
- The threshold is percentage-based at **60% of the maximum possible score**, allowing it to scale automatically if the rubric is changed in the future.

Output of the Scanning Agent

For each criterion, the agent provides:

- The criterion name
- Score out of 20
- A brief note explaining the score using a specific observation from the document

The agent also provides:

- **Overall Score:** Sum of the three criterion scores
- **Summary:** A short explanation of the main factors affecting the document's score.

The scoring is performed with a deterministic model configuration (`temperature = 0`) and the returned overall score is independently recalculated from the individual criterion scores rather than relying on the model's arithmetic.